

# TIME SERIES MODELING

*Jonathan Balaban*

*DAT<sub>2</sub>*

---

## TIME SERIES MODELING

---

# LEARNING OBJECTIVES

- ▶ Model and predict from time series data using AR, ARMA, or ARIMA models
- ▶ Specifically, coding these models in `statsmodels`

---

# PRE-WORK REVIEW

---

- ▶ Prior definition and Python functions for moving averages and autocorrelation
- ▶ Prior exposure to linear regression with discussion of coefficients and residuals
- ▶ `pip install statsmodels` (should be included with Anaconda)

---

# TIME SERIES MODELING

---

- ▶ In the last class, we focused on exploring time series data and common statistics for time series analysis.
- ▶ In this class, we will advance those techniques to show how to predict or forecast forward from time series data.
- ▶ With a sequence of values (a time series), we will use the techniques in this class to predict a future value.

---

# TIME SERIES MODELING

---

- ▶ There are many times when you may want to use a series of values to predict a future value.
  - ▶ The number of sales in a future month
  - ▶ Anticipated website traffic when buying a server
  - ▶ Financial forecasting
  - ▶ The number of visitors to your store during the holidays


## INTRODUCTION

---

# WHAT ARE TIME SERIES MODELS?

# WHAT ARE TIME SERIES MODELS?

- ▶ Time series models are models that will be used to predict a future value in the time series.
- ▶ **Like** other predictive models, we will use prior history to predict the future.
- ▶ **Unlike** previous models, we will use the earlier in time *outcome* variables as *inputs* for predictions.



---

# WHAT ARE TIME SERIES MODELS?

---

- ▶ **Like** previous modeling exercises, we will have to evaluate the different types of models to ensure we have chosen the best one.
- ▶ We will want to evaluate on a held-out set or test data to ensure our model performs well on unseen data.



---

# WHAT ARE TIME SERIES MODELS?

---

- ▶ **Unlike** previous modeling exercises, we won't be able to use standard cross-validation for evaluation.
- ▶ Since there is a time component to our data, we cannot choose training and test examples at random.
- ▶ Suppose we did select a random 80% sample of data points for training and a random 20% for testing. What could go wrong?

---


# WHAT ARE TIME SERIES MODELS?

---

- ▶ The training dataset would likely contain data from *before* AND *after* a test dataset.
- ▶ This would not be possible in real life (you can't use future, unseen data points when building your model). Therefore, it's not a valid test of how our model would perform in practice.

# WHAT ARE TIME SERIES MODELS?

- Instead, we will exclusively train on values earlier (in time) in our data and test our model on values at the end of the data period.



---

# PROPERTIES FOR TIME-SERIES PREDICTION

---

- ▶ A *moving average* is an average of  $p$  surrounding data points in time.

$$F_t = \frac{1}{p} \sum_{k=t-p+1}^{t-p+1} Y_k$$

... to  $t - p + 1$ . This includes  $t, t + 1, t + 2, \dots, t-p, t-p+1$ .

Divide by the  $p$  surrounding data points

Get the  $p$  points from  $t$  ...

---

# PROPERTIES FOR TIME-SERIES PREDICTION

---

- ▶ *Autocorrelation* is how correlated a variable is with itself. Specifically, how related are variables earlier in time with variables later in time.

$$r_k = \frac{\sum_{t=k+1}^n (y_t - \bar{y})(y_{t-k} - \bar{y})}{\sum_{t=1}^n (y_t - \bar{y})^2}$$

- ▶ We fix a *lag*,  $k$ , which is how many time points earlier we should use to compute the correlation.

---

# PROPERTIES FOR TIME-SERIES PREDICTION

---

- ▶ We can use these values to assess how we plan to model our time series.
- ▶ Typically, for a high quality model, we require some autocorrelation in our data.
- ▶ We can compute autocorrelation at various lag values to determine how far back in time we need to go.

---


# PROPERTIES FOR TIME-SERIES PREDICTION

---


- ▶ Many models make an assumption of *stationarity*, assuming the mean and variance of our values is the *same* throughout.
- ▶ While the values (e.g. of sales) may shift up or down over time, the mean and variance of sales is constant (i.e. there aren't many dramatic swings up or down).
- ▶ These assumptions may not represent real world data; we must be aware of that when we are breaking the assumptions of our model.

# PROPERTIES FOR TIME-SERIES PREDICTION

**Stationary Time Series**



**Non-stationary Time Series**






---

# PROPERTIES FOR TIME-SERIES PREDICTION

---


- For example, typical stock or market performance is not stationary. In this plot of Dow Jones performance since 1986, the mean is clearly increasing over time.



# PROPERTIES FOR TIME-SERIES PREDICTION

- ▶ Below are simulated examples of non-stationary time series and why they might occur.

**Table 1 Non-stationary behavior**



---


# PROPERTIES FOR TIME-SERIES PREDICTION

---

- ▶ Often, if these assumptions don't hold, we can alter our data to make them true. Two common methods are *detrending* and *differencing*.
- ▶ *Detrending* would mean to remove any major trends in our data.
- ▶ We could do this in many ways, but the simplest is to fit a line to the trend and make a new series that is the difference between the line and the true series.

# PROPERTIES FOR TIME-SERIES PREDICTION


- For example, there is a clear upward (non-stationary) trend in google searches for “iphone”.



- If we fit a line to this data first, we can create a new series that is the difference between the true number of searches and the predicted searches.


# PROPERTIES FOR TIME-SERIES PREDICTION

- ▶ Below is an example where we look at US housing prices over time. Clearly, there is an upward trend, making the time series non-stationary (i.e.: the mean house price is increasing).




# PROPERTIES FOR TIME-SERIES PREDICTION

- We can fit a line that represents the trend. With our trend line, we can subtract the trend line value from the original value to get the bottom figure.



# PROPERTIES FOR TIME-SERIES PREDICTION

- The data now has a fixed mean and will be easier to model. This pattern is similar to mean-scaling our features in earlier models with `StandardScaler`.



---

# PROPERTIES FOR TIME-SERIES PREDICTION


---

- ▶ A simpler method is *differencing*. This is very closely related to the `diff` function we saw in the last class.
- ▶ Instead of predicting the series (again our non-stationary series), we can predict the difference between two consecutive values.




# PROPERTIES FOR TIME-SERIES PREDICTION

Box-Cox(Passengers)



Box-Cox(Passengers)

Differencing (Box-Cox(Passengers))



Box-Cox(Passengers)

---

# TIME SERIES MODELS

---

- ▶ In the rest of this lesson, we are going to build up to the **ARIMA** time series model.
- ▶ This model combines the ideas of differencing and:
  - ▶ **AR** - autoregressive models
  - ▶ **MA** - moving average models



---

# AR MODELS

---

- ▶ Autoregressive (AR) models are those that use data from previous time points to predict the next.
- ▶ This is very similar to previous regression models, except as input, we take the previous outcome.
- ▶ If we are attempting to predict weekly sales, we use the sales from a previous week as input.
- ▶  $AR(p)$ , where  $p$  indicates the number of previous time points to incorporate, with  $AR(1)$  being the most common.

# AR MODELS



---

# AR MODELS

---

- ▶ In an autoregressive model, similar to standard regression, we are learning regression coefficients for each of the  $p$  previous values. Therefore, we will learn  $p$  coefficients or  $\beta$  values.
- ▶ If we have a time series of sales per week,  $y_i$ , we can regress each  $y_i$  from the last  $p$  values.

$$y_i = \beta_0 + \beta_1 y_{i-1} + \beta_2 y_{i-2} + \dots + \beta_p y_{i-p} + \varepsilon$$

- ▶ As with standard regression, our model assumes that each outcome variable is a linear combination of the inputs and a random error term.

---

# AR MODELS

---


- ▶ For an AR(1) model, we will learn a single coefficient.
- ▶ This coefficient,  $\beta$ , will tell us the relationship between the previous value,  $Y_{t-1}$ , and the next value,  $Y_t$ .

$$Y_t = \beta \cdot Y_{t-1}$$

# AR MODELS

---


- ▶ A value  $> 1$  would indicate a growth over previous values. This would typically represent non-stationary data, since if we compound the increases, the values are continually increasing.



# AR MODELS

---

- Values between 1 and -1 represent increasing and decreasing patterns from previous patterns.





---

# AR MODELS

---

- ▶ As with other models, interpretation of the model becomes more complex as we add more factors.
- ▶ Going from AR(1) to AR(2) can add significant *multi-collinearity*.

---

# AR MODELS

---

- ▶ Recall that *autocorrelation* is the correlation of a value with its series *lagged* behind.
- ▶ A model with high correlation implies that the data is highly dependent on previous values and an autoregressive model would perform well.

---

# AR MODELS

---

- ▶ Autoregressive models are useful for learning falls or rises in our series.
- ▶ This will weight together the last few values to make a future prediction.
- ▶ Typically, this model type is useful for small-scale trends such as an increase in demand or change in tastes that will gradually increase or decrease the series.

---

# ACTIVITY: KNOWLEDGE CHECK

---

## ANSWER THE FOLLOWING QUESTIONS



### EXERCISE

1. If we observe an autocorrelation near 1 for lag 1, what do we expect the single coefficient in an AR(1) model to be?  $>1$ , between 0 and 1, or  $<1$ ?
2. What if we observe an autocorrelation of 0?

---

# MA MODELS

---

- ▶ **Moving average (MA) models**, as opposed to AR models, do not take the previous outputs (or values) as inputs. They take the previous error terms.
- ▶ We will attempt to predict the next value based on the overall average and how off our previous predictions were.

---

# MA MODELS

---

- ▶ This model is useful for handling specific or abrupt changes in a system.
- ▶ AR models slowly incorporate changes in the system by combining previous values; MA models use prior errors to quickly incorporate changes.
- ▶ This is useful for modeling a sudden occurrence - something going out of stock, or a sudden rise in popularity affecting sales.

---

# MA MODELS

---

- ▶ As in AR models, we have an order term,  $q$ , and we refer to our model as MA( $q$ ). The moving average model is dependent on the last  $q$  errors.
- ▶ If we have a time series of sales per week,  $y_i$ , we can regress each  $y_i$  from the last  $q$  error terms.

$$y_i = \text{mean} + \beta_1 \varepsilon_{i-1} + \beta_2 \varepsilon_{i-2} + \dots + \beta_q \varepsilon_{i-q}$$

- ▶ We include the mean of the time series (that's why it's called a moving average) as we assume the model takes the mean value of the series and randomly jumps around it.

---

# MA MODELS

---

- ▶ Of course, we don't have error terms when we start - where do they come from?
- ▶ This requires a more complex fitting procedure than we have seen previously.
- ▶ We need to iteratively fit a model (perhaps with random error terms), compute the errors and then refit, again and again.



---


# MA MODELS

---

- ▶ In this model, we learn  $q$  coefficients.
- ▶ In an MA(1) model, we learn one coefficient.
- ▶ This value indicates the impact of our previous error term on the next prediction.

# ARMA MODELS

- ▶ **ARMA** (pronounced 'R-mah') models combine the autoregressive and moving average models.
- ▶ An ARMA(p,q) model is simply a combination (sum) of an AR(p) model and MA(q) model.



---

# ARMA MODELS

---

- ▶ We specify two model settings,  $p$  and  $q$ , which correspond to combining an AR( $p$ ) model with an MA( $q$ ) model.
- ▶ Incorporating both models allows us to mix two types of effects.
  - ▶ AR models slowly incorporate changes in preferences, tastes, and patterns.
  - ▶ Moving average models base their prediction on the prior error, allowing to correct sudden changes based on random events - supply, popularity spikes, etc.

---

# ARIMA MODELS

---

- ▶ **ARIMA** (pronounced ‘uh-ri-mah’) is an **AutoRegressive Integrated Moving Average** model.
- ▶ In this model, we learn an ARMA(p,q) model to predict *the difference* of the series (as opposed to the value of the series).

---

# ARIMA MODELS

---

- ▶ Recall the pandas `diff` function. This computes the difference between two consecutive values.
- ▶ In an ARIMA model, we attempt to predict this difference instead of the actual values.

$$y_t - y_{t-1} = \text{ARIMA}(p,q)$$

- ▶ This handles the stationarity assumption we wanted for our data. Instead of detrending or differencing manually, the model does it for us!

---

# ARIMA MODELS

---

- ▶ An ARIMA model has three parameters and is specified ARIMA( $p$ ,  $d$ ,  $q$ ).
  - ▶  $p$  is the order of the autoregressive component
  - ▶  $q$  is the order of the moving average component
  - ▶  $d$  is the degree of differencing.
- ▶  $d$  was 1 in our prior example. For  $d=2$ , our model would be

$$\text{diff}(\text{diff}(y)) = (y_t - y_{t-1}) - (y_{t-1} - y_{t-2}) = \text{ARIMA}(p,q)$$

---

# ARIMA MODELS

---

- ▶ Compared to an ARMA model, ARIMA models do **not** rely on the underlying series being stationary.
- ▶ The differencing operation can *convert* the series to one that is stationary.
- ▶ Instead of attempting to predict values over time, our new series is the difference in values over time.
- ▶ Since ARIMA models include differencing, they can be used on a broader set of data without the assumption of a constant mean.

---

**DEMO**

---

# TIME SERIES MODELING IN STATSMODELS



---

# TIME SERIES MODELING IN STATSMODELS

---

- ▶ To explore time series models, we will continue to use the Rossmann sales data.
- ▶ This dataset has sales data for every Rossmann store for a 3-year period and indicators for holidays and basic store information.

---

# TIME SERIES MODELING IN STATSMODELS

---

- ▶ In the last class, we saw that we could plot the sales data at a particular store to identify how the sales changed over time.
- ▶ We also computed autocorrelation for the data at varying lag periods.
- ▶ This helps us identify if previous timepoints are predictive of future data and which time points are most important - the previous day, week, or month.

---

# ACTIVITY: KNOWLEDGE CHECK

---

## ANSWER THE FOLLOWING QUESTIONS



EXERCISE

1. Compute the autocorrelation of Sales in Store 1 for lag 1 and 2.
2. Will we be able to use a predictive model, particularly an autoregressive one?

---

# TIME SERIES MODELING IN STATSMODELS

---

```
store1_data.Sales.autocorr(lag=1) # -0.12  
store1_data.Sales.autocorr(lag=2) # -0.03
```

- We do see some minimal correlation in time, implying an AR model can be useful. An easier way to diagnose this may be to plot many autocorrelations at once.

---

# TIME SERIES MODELING IN STATSMODELS

---

```
%matplotlib inline  
from pandas.tools.plotting import autocorrelation_plot  
  
autocorrelation_plot(store1_data.Sales)
```

- ▶ This shows a typical pattern of an autocorrelation plot, that it should decrease to 0 as lag increases. However, it's hard to observe exactly what the values are.

---

# TIME SERIES MODELING IN STATSMODELS

---

- ▶ In this class, we will use `statsmodels` to code AR, MA, ARMA, and ARIMA models.
- ▶ `statsmodels` provides a nice summary utility to help us diagnose models.

---

# TIME SERIES MODELING IN STATSMODELS

---

- ▶ statsmodels also has a better autocorrelation plot that allows us to look at fixed number of lag values.

```
from statsmodels.graphics.tsaplots import plot_acf  
plot_acf(store1_data.Sales, lags=10)
```

- ▶ Here we observe autocorrelation at 10 lag values. 1 and 2 are what we saw before.
- ▶ This implies a small but limited impact based on the last few values. An autoregressive model might be useful.

---

# TIME SERIES MODELING IN STATSMODELS

---

- ▶ We also see a larger spike at 7 (the seventh day in the week).
- ▶ If we observed a handful of random distributed spikes, a moving average model would be useful.

```
plot_acf(store1_data.Sales, lags=25)
```

- ▶ We can expand the window to 25 days to see that the random spikes occur regularly at 7 days. What does this mean?



---

# AR, MA, AND ARMA MODELS IN STATSMODELS

---

- ▶ To explore AR, MA, and ARMA models, we will use `sm.tsa.ARMA`.
- ▶ Remember, an ARMA model is a combination of autoregressive and moving average models.
- ▶ We can train an AR model by turning off the MA component ( $q=0$ ).

```
from statsmodels.tsa.arima_model import ARMA
```

```
store1_sales_data = store1_open_data[['Sales']].astype(float)  
model = ARMA(store1_sales_data, (1, 0)).fit()  
model.summary()
```

---

## AR, MA, AND ARMA MODELS IN STATSMODELS

---

- ▶ By passing  $(1, 0)$  in the second argument, we are fitting an ARMA model with  $p=1$ ,  $q=0$ . This is the same as an AR(1) model.
- ▶ In this AR(1) model, we learn an intercept (or base sales) value.
- ▶ Additionally, we learn a coefficient that tells us how to include the latest sales value.
- ▶ In this case, we add an intercept of  $\sim 4700$  to  $0.68$  times the previous month's sales. Note that the coefficient is not equal to the lag 1 autocorrelation. This implies the data is **not** stationary.

---

# AR, MA, AND ARMA MODELS IN STATSMODELS

---

- ▶ We can learn an AR(2) model, which regresses each sales value on the last two.

```
model = ARMA(store1_sales_data, (2, 0)).fit()  
model.summary()
```

- ▶ In this case, we learn two coefficients, which tell us the effect of the last two sales values on the current sales.
- ▶ While this model may perform better, it may be more difficult to interpret.

---

# AR, MA, AND ARMA MODELS IN STATSMODELS

---

- ▶ Residuals are the errors of the model or how off our predictions are.
- ▶ Ideally, we want randomly distributed errors that are small.
- ▶ If the errors are large, our model does not perform well.
- ▶ If the errors have a pattern, particularly over time, we may have overlooked something in the model or have periods of time that are different than the rest of the dataset.

---

# AR, MA, AND ARMA MODELS IN STATSMODELS

---

- ▶ We can use `statsmodels` to plot the residuals.

```
model.resid.plot()
```

- ▶ Here we see large spikes at the end of each year, indicating that our model does not account for the holiday spikes.
- ▶ Our model considers a short period of time, so it does not take into account the longer seasonal pattern.

---

# AR, MA, AND ARMA MODELS IN STATSMODELS

---

- ▶ We can also plot the autocorrelations of the residuals. In an ideal world, these would all be near 0 and appear random.

```
plot_acf(model.resid, lags=50)
```

- ▶ This plot shows a problem: the errors are increasing and decreasing every week in a clear pattern.
- ▶ We may need to expand our model.

---

# AR, MA, AND ARMA MODELS IN STATSMODELS

---

- To expand this AR model to an ARMA model, we can include the moving average component as well.

```
model = ARMA(store1_sales_data, (1, 1)).fit()  
model.summary()
```

- Now we learn two coefficients, one for the AR(1) component and one for the MA(1) component.

---

## AR, MA, AND ARMA MODELS IN STATSMODELS

---

- ▶ Remember that this is an AR(1) + MA(1) model. The AR coefficient represents dependency on the last value and the MA component represents any spikes independent of the last value.
- ▶ The coefficients here are 0.69 for the AR component and -0.03 for the MA component.
- ▶ The AR coefficient is the same as before (decreasing values).
- ▶ The MA component is fairly small (which we should have expected from the autocorrelation plots).



---

# ARIMA MODELS IN STATSMODELS

---

- ▶ We can also use `statsmodels` to fit ARIMA models. Let's start by using `ARIMA(1, 0, 1)` to fit an ARMA(1, 1) model.

```
from statsmodels.tsa.arima_model import ARIMA
```

```
model = ARIMA(store1_sales_data, (1, 0, 1)).fit()  
model.summary()
```

- ▶ We can see that this model is the same as our previous ARMA model.

---

# ARIMA MODELS IN STATSMODELS

---

- ▶ We can also fit a true ARIMA model to predict the difference of the series.

```
model = ARIMA(store1_sales_data, (1, 1, 1)).fit()  
model.summary()
```

- ▶ We can remove the MA component since it does not appear to be useful.

```
model = ARIMA(store1_sales_data, (1, 1, 0)).fit()  
model.summary()
```

- ▶ We now have an AR(1) model on the differenced series with a coefficient of -0.18.

---

# ARIMA MODELS IN STATSMODELS

---

- ▶ We can compute the lag 1 autocorrelation of the differenced series and see if they match.

```
store1_sales_data.Sales.diff(1).autocorr(1) #-0.181
```

- ▶ We can also plot it to see the difference.

```
store1_sales_data.Sales.diff(1).plot()
```

- ▶ They match. Note that this is generally true, but the variance is NOT constant. It's mostly the same throughout the series except around the holidays.

---

# ARIMA MODELS IN STATSMODELS

---

- ▶ With our models, we can also plot our predictions against the true series using the `plot_predict` function.
- ▶ We can compare the last 50 days of true values against our predictions.

```
model.plot_predict(0, 50)
```

- ▶ The function takes two arguments, the start and end index of the dataframe to plot. Here, we are plotting the last 50 values.

---

# ARIMA MODELS IN STATSMODELS

---

- ▶ To plot earlier values with our predictions continuing where the true values stop, we can do the following.

```
import matplotlib.pyplot as plt
```

```
fig, ax = plt.subplots()  
ax = store1_sales_data['2014'].plot(ax=ax)
```

```
fig = model.plot_predict(0, 200, ax=ax, plot_insample=False)
```

- ▶ This plots true values in 2014 and our predictions 200 days out from 2014.

---

# ARIMA MODELS IN STATSMODELS

---

- ▶ The two previous problems remain:
  - ▶ Large errors around the holiday period
  - ▶ Errors with high autocorrelation

---

# ARIMA MODELS IN STATSMODELS

---

- ▶ We can adjust the AR component of the model to adjust for a piece of this. Let's increase the lag to 7.

```
model = ARIMA(store1_sales_data, (7, 1, 2)).fit()  
model.summary()
```

```
plot_acf(model.resid, lags=50)
```

- ▶ This removes some of the autocorrelation in the residuals but large discrepancies still exist.
- ▶ However, they exist where we are breaking our model assumptions.

---

# ARIMA MODELS IN STATSMODELS

---

- ▶ Increasing  $p$  increases the dependency on previous values further (longer lag). But our autocorrelation plots show this isn't necessary past a certain point.
- ▶ Increasing  $q$  increases the likelihood of an unexpected jump at a handful of points. The autocorrelation plots show this doesn't help past a certain point.
- ▶ Increasing  $d$  increases differencing, but  $d=1$  moves our data towards stationarity (other than a few points).  $d=2$  would imply an exponential trend which we don't have here.



---

# ARIMA MODELS IN STATSMODELS

---

- ▶ There are variants of ARIMA that will better handle the seasonal aspect of our data. This is referred to as Seasonal ARIMA.
- ▶ These models fit two ARIMA models, one on the current frequency (daily in our example) and another on the seasonal frequency (maybe monthly or yearly patterns).
- ▶ Additionally, issues with seasonality could be handled by preprocessing tricks such as detrending.

---

**INDEPENDENT PRACTICE**

---

# **WALMART SALES DATA**

---

# ACTIVITY: WALMART SALES DATA

---



## EXERCISE

### DIRECTIONS (50 minutes)

We will analyze the weekly sales data from Walmart over a two year period from 2010 to 2012. The data is separated by store and department, but we will focus on analyzing one store for simplicity.

To read in the data

```
import pandas as pd
import numpy as np
```

```
%matplotlib inline
```

```
data = pd.read_csv('lessons/lesson-16/assets/data/
train.csv')
data.set_index('Date', inplace=True)
data.head()
```

---

# ACTIVITY: WALMART SALES DATA

---

## DIRECTIONS

Complete the following tasks:

1. Filter the dataframe to Store 1 sales and aggregate over departments to compute the total sales per store.
2. Plot the `rolling_mean` for `Weekly_Sales`. What general trends do you observe?
3. Compute the 1, 2, 52 autocorrelations for `Weekly_Sales` and/or create an autocorrelation plot.
4. What does the autocorrelation plot say about the type of model you want to build?



EXERCISE

---

# ACTIVITY: WALMART SALES DATA

---

## DIRECTIONS



### EXERCISE

5. Split the weekly sales data in a training and test set - using 75% of the data for training.
6. Create an AR(1) model on the training data and compute the mean absolute error of the predictions.
7. Plot the residuals - where are their significant errors?
8. Compute and AR(2) model and an ARMA(2, 2) model - does this improve your mean absolute error on the held out set?
9. Finally, compute an ARIMA model to improve your prediction error - iterate on the p, q, and parameters comparing the model's performance..

---

# CONCLUSION

---

- ▶ Time-series models use previous values to predict future values, also known as forecasting.
- ▶ AR and MA model are simple models on previous values or previous errors respectively.
- ▶ ARMA combines these two types of models to account for both gradual shifts (due to AR models) and abrupt changes (MA models).

---

# CONCLUSION

---

- ▶ ARIMA models train ARMA models on differenced data to account for non-stationary data.
- ▶ Note that none of these models may perform well for data that has more random variation.
- ▶ For something like iPhone sales (or searches), which may be sporadic with short periods of increase, these models may not work well.

---

**COURSE**

---

**BEFORE NEXT  
CLASS**

**Final: Part 4**



---

# LESSON

---

Q & A

## LESSON

---

# EXIT TICKET

**DON'T FORGET TO FILL OUT YOUR EXIT  
TICKET**